


Project: File System

Textbook: pages 501-506

Lubomir Bic

Assignment

- Design and implement a simple file system using ldisk to emulate disk
- Overall organization



- Input
 - `cr foo`
 - `op foo`
 - `wr 1 y 10`
 - `sk 1 0`
 - `rd 1 3`
- Output
 - `file foo created`
 - `file foo opened, index=1`
 - `10 bytes written`
 - `current position is 0`
 - `3 bytes read: yyy`

I/O System

- I/O system presents disk as a linear sequence of blocks:
 - We will refer to the logical disk as *ldisk[L][B]*
 - L is the number of logical blocks on ldisk
 - B is the block length (in bytes)
 - It can be implemented as a byte array or integer array
 - Type casting or conversion is necessary
- I/O system interface – provided by your driver:
`read_block(int i, char *p)`
`write_block(int i, char *p)`
- Each command reads or writes an entire block (B bytes)
- FS can access the emulated disk using only these functions
(no direct access to ldisk is allowed)

File System -- User Interface

- `create(symbolic_file_name)`: return ok/error
- `destroy(symbolic_file_name)` : return ok/error
- `open(symbolic_file_name)`: return index/error
- `close(index)`: return ok/error
- `read(index, mem_area, count)`: return #bytes read/error
- `write(index, mem_area, count)`: return #bytes written/error
- `lseek(index, pos)` : return ok/error
- `directory`: return list of files/error
- `init/save`: create or restore ldisk/save ldisk

Review of concepts

- directory structure
 - tree, DAG, graph, symbolic links, path names
 - this project: single flat list of all files (=one special file)
- organization of entries within directory
 - array of slots, linked list, hash table, B-tree
 - this project: unsorted array of fixed-size slots
- each directory entry contains
 - all descriptive info, parts of info, name only
 - this project: symbolic name plus index of descriptor

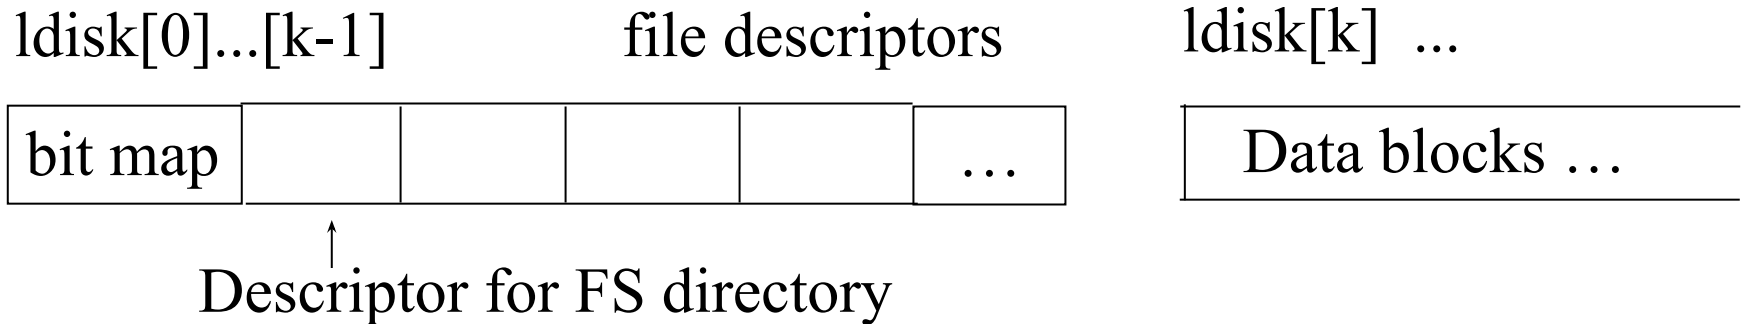
Review of concepts

- file descriptor contents
 - owner, file type, protection, length, disk map, access times
 - this project: length (bytes), disk map
- disk map (physical organization)
 - contiguous, linked list, indexed, multi-level
 - this project:
 - flat index (fixed list of max 3 disk blocks)
 - 1-level incremental (for teams)


Review of concepts

- location of file descriptors
 - dedicated portion of disk, special files, in directories
 - this project: first k disk blocks ($ldisk[0]..[k-1]$)
- free storage management
 - linked lists, bit map
 - this project: bit map

Organization of the file system



Each descriptor



- teams: additional task: 1-level incremental index
 - last entry points to another file descriptor (2+4 block)



The Directory

- only one directory (root)
- regular file, i.e., use regular file operations:
 - read, write, lseek
 - but the directory is always open (OFT[0])
- described by the first descriptor
- contains array of entries:
 - symbolic file name (characters)
 - index of the descriptor (integer)

Create a file

- find a free file descriptor
- find a free directory entry
- fill both entries

i




Destroy a file

- search directory to find file descriptor
- remove directory entry
- update bit map (if file was not empty)
- free file descriptor
- return status

Open a file

OFT:



- search directory to find index of file descriptor (i)
- allocate a free OFT entry (reuse deleted entries)
- fill in current position (0) and file descriptor index (i)
- read block 0 of file into the r/w buffer (read-ahead)
- return OFT index (j) (or return error)
- consider adding a file length field (to simplify checking)

Close a file

- write buffer to disk
- update file length in descriptor
- free OFT entry
- return status

Read a file

- compute position in the r/w buffer
- copy from buffer to memory until
 - desired count or end of file is reached:
update current pos, return status
 - end of buffer is reached
 - write the buffer to disk
 - read the next block
 - continue copying

Write a file

- write into buffer
- when full, write buffer to disk
 - if block does not exist (file is expanding):
 - allocate new block
 - update file descriptor
 - update bit map
- update file length in descriptor

Seek in a file

- if the new position is not within the current block
 - write the buffer to disk
 - read the new block
- set the current position to the new position
- return status

List the directory

- read directory file
- for each entry:
 - find file descriptor
 - print file name and file length

Presentation shell

- develop presentation shell:
 - repeatedly accept command (e.g. `cr abc`)
 - invoke corresponding FS function (e.g. `create(abc)`)
 - display status/data on screen
(e.g. `file abc created or error`)
- project will be tested using an input file and it must produce an output file

Shell commands and Output

- `cr <name>`
 - **Output:** `file <name> created`
- `de <name>`
 - **Output:** `file <name> destroyed`
- `op <name>`
 - **Output:** `file <name> opened, index=<index>`
- `cl <index>`
 - **Output:** `file <name> closed`
- `rd <index> <count>`
 - **Output:** `<count> bytes read: <xx...x>`
- `wr <index> <char> <count>`
 - **Output:** `<count> bytes written`

Shell commands and Output

- `sk <index> <pos>`
 - **Output:** current position is `<pos>`
- `dr`
 - **Output:** `file0 <len0>, ..., fileN <lenN>`
- `in <no_cyl> <no_surf> <no_sect> <sect_len>`
`<disk_cont>`
 - `disk_cont` is a text file; it holds copy of `ldisk`
 - If file does not exist, output: `disk initialized`
 - If file does exist, output: `disk restored`
- `sv <disk_cont>`
 - **Output:** `disk saved`
- If any command fails, output: `error`

Sample Interaction

- Input

```
in 4 2 8 64 dsk.txt
cr foo
op foo
wr 1 x 60
wr 1 y 10
sk 1 55
rd 1 10
dr
sv dsk.txt
in 4 2 8 64 dsk.txt
op foo
rd 1 3
cr foo
```

- Output

```
disk initialized
file foo created
file foo opened, index=1
60 bytes written
10 bytes written
current position is 55
10 bytes read: xxxxxyyyyyy
foo 70
disk saved
disk restored
file foo opened, index=1
3 bytes read: xxx
error
```

Handling the Bit Map (pg 217)

- determine BM size (# of bits needed = # of ldisk blocks)
- represent bit map as an array of int (32 bits each):
BM [n]
- **How to set, reset, and search for bits in BM?**
- prepare a mask array: MASK[16]
 - diagonal contains “1”, all other fields are “0”
 - use bit operations (bitwise or/and) to manipulate bits

Handling the Bit Map

- MASK

0	10...
1	010...
2	0010...
3	00010...
...	...
16	0 ... 01

- to set bit i of $BM[j]$ to “1”:

$$BM[j] = BM[j] \mid MASK[i]$$

Handling the Bit Map

- how to create MASK?

MASK[0] = 0x8000 (1000 0000 0000 0000)

MASK[1] = 0x4000 (0100 0000 0000 0000)

MASK[2] = 0x2000 (0010 0000 0000 0000)

MASK[3] = 0x1000 (0001 0000 0000 0000)

MASK[4] = 0x0800 (0000 1000 0000 0000)

...

MASK[15] = 0x0001 (0000 0000 0000 0001)

- another approach:

MASK[15] = 1;

MASK[i] = MASK[i+1] <<

Handling the Bit Map

- to set a bit to “0”:
 - create MASK2, where $\text{MASK2}[i] = \sim\text{MASK}[i]$
e.g., 0010 0000 0000 0000 \rightarrow 1101 1111 1111 1111
(use “ \sim ” operator or declare using hex constants like MASK)
- set bit i of $\text{BM}[j]$ to “0”:
$$\text{BM}[j] = \text{BM}[j] \& \text{MASK2}[i]$$

Handling the Bit Map

- to search for a bit equal to “0” in BM:

```
for (i=0; ...      /* search BM from the beginning
    for (j=0; ...  /* check each bit in BM[i] for “0”
        test = BM[i] & MASK[j])
        if (test == 0) then bit j of BM[i] is “0”; stop search
```

Assumptions for Testing

- Disk: 4 cylinders, 2 surfaces, 8 sectors/track, 64 Bytes/sector
 - sector = block = 64 B = 16 int
- What is k?
 - Bitmap: $4 * 2 * 8 = 64$ bits (2 integers)
 - Descriptor: 4 integers (file length plus 3 block #s)
- How many descriptors (files) can we have?
 - How many fit into directory?
 - Each dir entry: 2 int
 - File name: maximum 4 chars, no extension (=1 int)
 - Descriptor index: 1 integer
 - dir = 3 blocks = $3 * 64 \text{ B} = 48 \text{ int} = 24 \text{ entries} (= 24 \text{ files})$
 - 24 descriptors = 96 int = 6 blocks
- k=7: 1 block for bitmap, 6 blocks for descriptors

Assumptions for Testing

- Number of open files: 3 plus the directory, i.e., the size of the OFT is exactly 4 entries.
- Directory should be opened automatically when program starts (index=0)
- It should close automatically with sv command
- All other files should also close at that time

Summary of tasks

- design and implement I/O system (ldisk plus read/write ops)
- design and implement FS using ldisk
- develop test/presentation shell
- teams:
 - error checks on all commands
 - additional task (pg 506): 1-level expanding index
- submit documentation
- schedule testing